

```
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
import datetime as dt
```

```
df = pd.read_csv("uber.csv")
df.head()
```

```
-----
-----
FileNotFoundError                                Traceback (most recent call
last)
Cell In[3], line 1
----> 1 df = pd.read_csv("uber.csv")
      2 df.head()

File ~\anaconda3\Lib\site-packages\pandas\util\_decorators.py:211, in
deprecate_kwarg.<locals>._deprecate_kwarg.<locals>.wrapper(*args,
**kwargs)
    209     else:
    210         kwargs[new_arg_name] = new_arg_value
--> 211 return func(*args, **kwargs)

File ~\anaconda3\Lib\site-packages\pandas\util\_decorators.py:331, in
deprecate_nonkeyword_arguments.<locals>.decorate.<locals>.wrapper(*arg
s, **kwargs)
    325 if len(args) > num_allow_args:
    326     warnings.warn(
    327
msg.format(arguments=_format_argument_list(allow_args)),
    328         FutureWarning,
    329         stacklevel=find_stack_level(),
    330     )
--> 331 return func(*args, **kwargs)

File ~\anaconda3\Lib\site-packages\pandas\io\parsers\readers.py:950,
in read_csv(filepath_or_buffer, sep, delimiter, header, names,
index_col, usecols, squeeze, prefix, mangle_dupe_cols, dtype, engine,
converters, true_values, false_values, skipinitialspace, skiprows,
skipfooter, nrows, na_values, keep_default_na, na_filter, verbose,
skip_blank_lines, parse_dates, infer_datetime_format, keep_date_col,
date_parser, dayfirst, cache_dates, iterator, chunksize, compression,
thousands, decimal, lineterminator, quotechar, quoting, doublequote,
escapechar, comment, encoding, encoding_errors, dialect,
error_bad_lines, warn_bad_lines, on_bad_lines, delim_whitespace,
low_memory, memory_map, float_precision, storage_options)
    935 kwds_defaults = _refine_defaults_read(
    936     dialect,
    937     delimiter,
```

```

(...)
946     defaults={"delimiter": ",",}
947 )
948 kwds.update(kwds_defaults)
--> 950 return _read(filepath_or_buffer, kwds)

File ~\anaconda3\Lib\site-packages\pandas\io\parsers\readers.py:605,
in _read(filepath_or_buffer, kwds)
    602 _validate_names(kwds.get("names", None))
    604 # Create the parser.
--> 605 parser = TextFileReader(filepath_or_buffer, **kwds)
    607 if chunksize or iterator:
    608     return parser

File ~\anaconda3\Lib\site-packages\pandas\io\parsers\readers.py:1442,
in TextFileReader.__init__(self, f, engine, **kwds)
    1439     self.options["has_index_names"] = kwds["has_index_names"]
    1441 self.handles: IOHandles | None = None
-> 1442 self._engine = self._make_engine(f, self.engine)

File ~\anaconda3\Lib\site-packages\pandas\io\parsers\readers.py:1735,
in TextFileReader._make_engine(self, f, engine)
    1733     if "b" not in mode:
    1734         mode += "b"
-> 1735 self.handles = get_handle(
    1736     f,
    1737     mode,
    1738     encoding=self.options.get("encoding", None),
    1739     compression=self.options.get("compression", None),
    1740     memory_map=self.options.get("memory_map", False),
    1741     is_text=is_text,
    1742     errors=self.options.get("encoding_errors", "strict"),
    1743     storage_options=self.options.get("storage_options", None),
    1744 )
    1745 assert self.handles is not None
    1746 f = self.handles.handle

File ~\anaconda3\Lib\site-packages\pandas\io\common.py:856, in
get_handle(path_or_buf, mode, encoding, compression, memory_map,
is_text, errors, storage_options)
    851 elif isinstance(handle, str):
    852     # Check whether the filename is to be opened in binary
mode.
    853     # Binary mode does not support 'encoding' and 'newline'.
    854     if ioargs.encoding and "b" not in ioargs.mode:
    855         # Encoding
--> 856         handle = open(
    857             handle,
    858             ioargs.mode,
    859             encoding=ioargs.encoding,

```

```

860         errors=errors,
861         newline="",
862     )
863     else:
864         # Binary mode
865         handle = open(handle, ioargs.mode)

```

FileNotFoundError: [Errno 2] No such file or directory: 'uber.csv'

```
df.drop(columns=['Unnamed: 0', 'key'], inplace=True)
```

```
df.info()
```

```

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 200000 entries, 0 to 199999
Data columns (total 7 columns):

```

#	Column	Non-Null Count	Dtype
0	fare_amount	200000 non-null	float64
1	pickup_datetime	200000 non-null	object
2	pickup_longitude	200000 non-null	float64
3	pickup_latitude	200000 non-null	float64
4	dropoff_longitude	199999 non-null	float64
5	dropoff_latitude	199999 non-null	float64
6	passenger_count	200000 non-null	int64

```

dtypes: float64(5), int64(1), object(1)
memory usage: 10.7+ MB

```

```
df.dropna(how='any', inplace=True)
```

```
df.isnull().sum()
```

```



fare_amount      0
pickup_datetime  0
pickup_longitude  0
pickup_latitude  0
dropoff_longitude 0
dropoff_latitude 0
passenger_count  0
dtype: int64



```



```

for col in df.select_dtypes(exclude=['object']):
    plt.figure()
    sns.boxplot(data=df, x=col)

```







```

df = df[
    (df.pickup_latitude > -90) & (df.pickup_latitude < 90) &
    (df.dropoff_latitude > -90) & (df.dropoff_latitude < 90) &
    (df.pickup_longitude > -180) & (df.pickup_longitude < 180) &
    (df.dropoff_longitude > -180) & (df.dropoff_longitude < 180) &
    (df.fare_amount > 0) & (df.passenger_count > 0) &
    (df.passenger_count < 50)
]

from math import cos, asin, sqrt, pi
import numpy as np

def distance(lat_1,lon_1,lat_2,lon_2):
#     lat1 = row.pickup_latitude
#     lon1 = row.pickup_longitude
#     lat2 = row.dropoff_latitude
#     lon2 = row.dropoff_longitude
    lon_1, lon_2, lat_1, lat_2 = map(np.radians, [lon_1, lon_2, lat_1,
lat_2]) #Degrees to Radians

    diff_lon = lon_2 - lon_1
    diff_lat = lat_2 - lat_1

    km = 2 * 6371 * np.arcsin(np.sqrt(np.sin(diff_lat/2.0)**2 +
np.cos(lat_1) * np.cos(lat_2) * np.sin(diff_lon/2.0)**2))

    return km

temp =
distance(df['pickup_latitude'],df['pickup_longitude'],df['dropoff_latitude'],df['dropoff_longitude'])
temp.head()

0      1.683323
1      2.457590
2      5.036377
3      1.661683
4      4.475450
dtype: float64

df_new = df.copy()
df_new['Distance'] = temp
df = df_new
df.head()

   fare_amount      pickup_datetime  pickup_longitude
pickup_latitude \
0           7.5  2015-05-07 19:52:06 UTC          -73.999817

```

```


40.738354
1      7.7  2009-07-17 20:04:56 UTC      -73.994355
40.728225
2     12.9  2009-08-24 21:45:00 UTC      -74.005043
40.740770
3      5.3  2009-06-26 08:22:21 UTC      -73.976124
40.790844
4     16.0  2014-08-28 17:47:00 UTC      -73.925023
40.744085

```

	dropoff_longitude	dropoff_latitude	passenger_count	Distance
0	-73.999512	40.723217	1	1.683323
1	-73.994710	40.750325	1	2.457590
2	-73.962565	40.772647	1	5.036377
3	-73.965316	40.803349	3	1.661683
4	-73.973082	40.761247	5	4.475450

```
sns.boxplot(data=df,x='Distance')
```

```
<Axes: xlabel='Distance'>
```



```
df = df[(df['Distance'] < 200) & (df['Distance'] > 0)]
```



```
df['pickup_datetime'] = pd.to_datetime(df['pickup_datetime'])
```

C:\Users\prajw\AppData\Local\Temp\ipykernel_26664\200932308.py:1:

SettingWithCopyWarning:

A value is trying to be set on a copy of a slice from a DataFrame.
Try using .loc[row_indexer,col_indexer] = value instead

See the caveats in the documentation:

https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy

```
df['pickup_datetime'] = pd.to_datetime(df['pickup_datetime'])
```

```
df['week_day'] = df['pickup_datetime'].dt.day_name()
```

```
df['Year'] = df['pickup_datetime'].dt.year
```

```
df['Month'] = df['pickup_datetime'].dt.month
```

```
df['Hour'] = df['pickup_datetime'].dt.hour
```

C:\Users\prajw\AppData\Local\Temp\ipykernel_26664\2592915223.py:1:

SettingWithCopyWarning:

A value is trying to be set on a copy of a slice from a DataFrame.
Try using .loc[row_indexer,col_indexer] = value instead

See the caveats in the documentation:

https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy

```
df['week_day'] = df['pickup_datetime'].dt.day_name()
```

C:\Users\prajw\AppData\Local\Temp\ipykernel_26664\2592915223.py:2:

SettingWithCopyWarning:

A value is trying to be set on a copy of a slice from a DataFrame.
Try using .loc[row_indexer,col_indexer] = value instead

See the caveats in the documentation:

https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy

```
df['Year'] = df['pickup_datetime'].dt.year
```

C:\Users\prajw\AppData\Local\Temp\ipykernel_26664\2592915223.py:3:

SettingWithCopyWarning:

A value is trying to be set on a copy of a slice from a DataFrame.
Try using .loc[row_indexer,col_indexer] = value instead

See the caveats in the documentation:

https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy

```
df['Month'] = df['pickup_datetime'].dt.month
```

C:\Users\prajw\AppData\Local\Temp\ipykernel_26664\2592915223.py:4:

SettingWithCopyWarning:

A value is trying to be set on a copy of a slice from a DataFrame.
Try using .loc[row_indexer,col_indexer] = value instead

See the caveats in the documentation:

https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy
df['Hour'] = df['pickup_datetime'].dt.hour

```
df.drop(columns=['pickup_datetime','pickup_latitude','pickup_longitude',  
'dropoff_latitude','dropoff_longitude'],inplace=True)
```

C:\Users\prajw\AppData\Local\Temp\ipykernel_26664\3782303944.py:1:
SettingWithCopyWarning:
A value is trying to be set on a copy of a slice from a DataFrame

See the caveats in the documentation:

https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy

```
df.drop(columns=['pickup_datetime','pickup_latitude','pickup_longitude',  
'dropoff_latitude','dropoff_longitude'],inplace=True)
```

```
df.head()
```

	fare_amount	passenger_count	Distance	week_day	Year	Month	Hour
0	7.5	1	1.683323	Thursday	2015	5	19
1	7.7	1	2.457590	Friday	2009	7	20
2	12.9	1	5.036377	Monday	2009	8	21
3	5.3	3	1.661683	Friday	2009	6	8
4	16.0	5	4.475450	Thursday	2014	8	17

```
temp = df.copy()
```

```
def convert_week_day(day):  
    if day in ['Monday','Tuesday','Wednesday','Thursday']:  
        return 0 # Weekday  
    return 1 # Weekend
```

```
def convert_hour(hour):  
    if 5 <= hour <= 12:  
        return 1  
    elif 12 < hour <= 17:  
        return 2  
    elif 17 < hour < 24:  
        return 3  
    return 0
```

```
df['week_day'] = temp['week_day'].apply(convert_week_day)  
df['Hour'] = temp['Hour'].apply(convert_hour)  
df.head()
```

C:\Users\prajw\AppData\Local\Temp\ipykernel_26664\3260682206.py:17:
SettingWithCopyWarning:
A value is trying to be set on a copy of a slice from a DataFrame.

Try using `.loc[row_indexer,col_indexer] = value` instead

See the caveats in the documentation:

https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy

```
df['week_day'] = temp['week_day'].apply(convert_week_day)
```

C:\Users\prajw\AppData\Local\Temp\ipykernel_26664\3260682206.py:18:

SettingWithCopyWarning:

A value is trying to be set on a copy of a slice from a DataFrame.

Try using `.loc[row_indexer,col_indexer] = value` instead

See the caveats in the documentation:

https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy

```
df['Hour'] = temp['Hour'].apply(convert_hour)
```


	fare_amount	passenger_count	Distance	week_day	Year	Month	Hour
0	7.5	1	1.683323	0	2015	5	3
1	7.7	1	2.457590	1	2009	7	3
2	12.9	1	5.036377	0	2009	8	3
3	5.3	3	1.661683	1	2009	6	1
4	16.0	5	4.475450	0	2014	8	2

`df.corr()`

	fare_amount	passenger_count	Distance	week_day
Year \				
fare_amount	1.000000	0.011884	0.778667	0.002305
0.120430				
passenger_count	0.011884	1.000000	0.005112	0.035882
0.005339				
Distance	0.778667	0.005112	1.000000	0.014518
0.018617				
week_day	0.002305	0.035882	0.014518	1.000000
0.006910				
Year	0.120430	0.005339	0.018617	0.006910
1.000000				
Month	0.024120	0.008818	0.007373	-0.007328
0.115182				
Hour	-0.021078	0.013572	-0.022691	-0.078129
0.001131				

	Month	Hour
fare_amount	0.024120	-0.021078
passenger_count	0.008818	0.013572
Distance	0.007373	-0.022691
week_day	-0.007328	-0.078129
Year	-0.115182	0.001131
Month	1.000000	-0.005410
Hour	-0.005410	1.000000

```
sns.scatterplot(y=df['fare_amount'],x=df['Distance'])  
<Axes: xlabel='Distance', ylabel='fare_amount'>
```



```
from sklearn.preprocessing import StandardScaler  
x = df[['Distance']].values  
y = df['fare_amount'].values.reshape(-1,1)  
  
from sklearn.model_selection import train_test_split  
x_train, x_test, y_train,y_test =  
train_test_split(x,y,random_state=10)  
  
std_x = StandardScaler()  
x_train = std_x.fit_transform(x_train)  
x_test = std_x.transform(x_test)  
  
std_y = StandardScaler()  
y_train = std_y.fit_transform(y_train)  
y_test = std_y.transform(y_test)
```

```

from sklearn.metrics import mean_squared_error, r2_score,
mean_absolute_error
def fit_predict(model):
    model.fit(x_train, y_train.ravel())
    y_pred = model.predict(x_test)
    r_squared = r2_score(y_test, y_pred)
    RMSE = mean_squared_error(y_test, y_pred, squared=False)
    MAE = mean_absolute_error(y_test, y_pred)
    print('R-squared: ', r_squared)
    print('RMSE: ', RMSE)
    print("MAE: ", MAE)

from sklearn.linear_model import LinearRegression

fit_predict(LinearRegression())

R-squared: 0.604116792084117
RMSE: 0.6290054895695945
MAE: 0.2755232959095982

from sklearn.ensemble import RandomForestRegressor
fit_predict(RandomForestRegressor())

R-squared: 0.6526968425632438
RMSE: 0.589149157739537
MAE: 0.2919005686906083

model = RandomForestRegressor()
model.fit(x_train, y_train)


C:\Users\prajw\anaconda3\Lib\site-packages\sklearn\base.py:1151:
DataConversionWarning: A column-vector y was passed when a 1d array
was expected. Please change the shape of y to (n_samples,), for
example using ravel().
    return fit_method(estimator, *args, **kwargs)

RandomForestRegressor()

y_pred = model.predict(x_test)

sns.regplot(x=y_test, y=y_pred, color="red", line_kws={"color" :
"blue"})
plt.show()

```



```
model = LinearRegression()  
model.fit(x_train, y_train)  
scaler = StandardScaler()
```

```
-----  
-----  
NameError                                Traceback (most recent call  
last)  
Cell In[1], line 1  
----> 1 model = LinearRegression()  
      2 model.fit(x_train, y_train)  
      3 scaler = StandardScaler()
```

NameError: name 'LinearRegression' is not defined

```
y_pred = model.predict(x_test)
```

```
sns.regplot(x=y_test, y=y_pred, color="red", line_kws={"color" :  
"blue"})  
plt.show()
```